


Stream Processing with Kafka & KSQL

Nick Dearden
April 2020



Today's Agenda



10:00 - 10:45 AM

Streams Processing/KSQL Overview

Nick Dearden, Director Advanced Technologies, Confluent

10:45 AM - 12:15 PM

Interactive Streams Lab

Nick Dearden, Director Advanced Technologies, Confluent
Brian Likosar, Solutions Engineer, Confluent

12:15 - 12:30 PM

Q&A and Next Steps

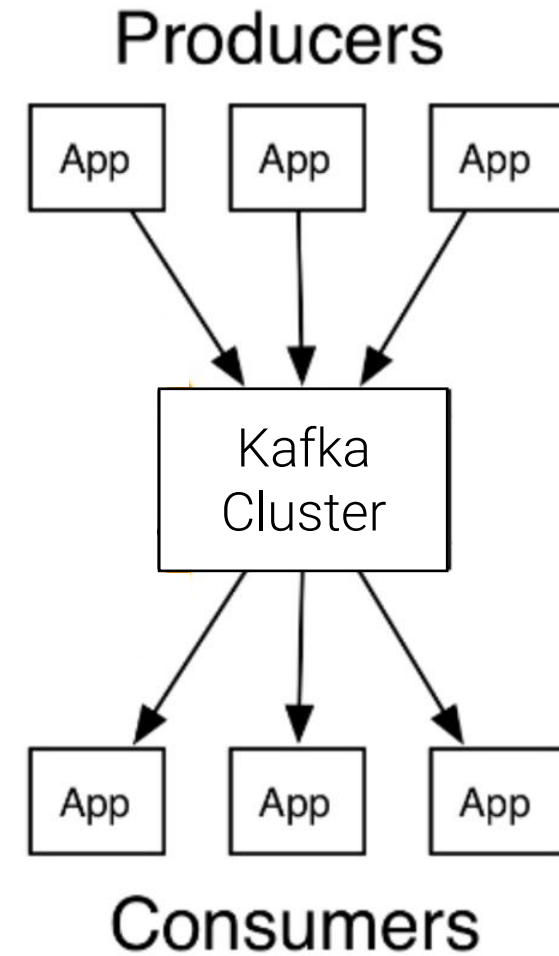

Open Discussion

Workshop Tips & Help:

1. Your username to log in is:
first initial last name
[all one word, lowercase]
2. Check the **'Chat' window** during the session for instructions
[icon located at the bottom of the Zoom toolbar]
3. For any technical issues, click the **'Raise Hand' button** or post in the **'Chat' window**
[a Confluent team member will assist you]

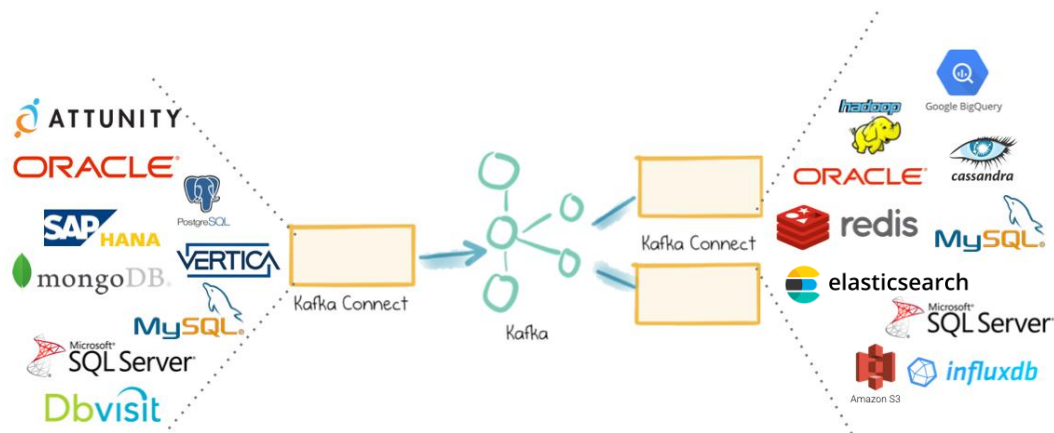
Kafka

A **Distributed Commit Log**. Publish and subscribe to streams of records. Highly scalable, high throughput. Supports transactions. Persisted data.




Kafka Connect API

Reliable and scalable integration of Kafka with other systems – no coding required.




Kafka Streams API

Write standard Java applications & microservices to process your data in real-time




What does a streaming platform do?

**Publish and
subscribe to streams
of data**




similar to a message
queue or enterprise
messaging system.

**Store streams
of data**



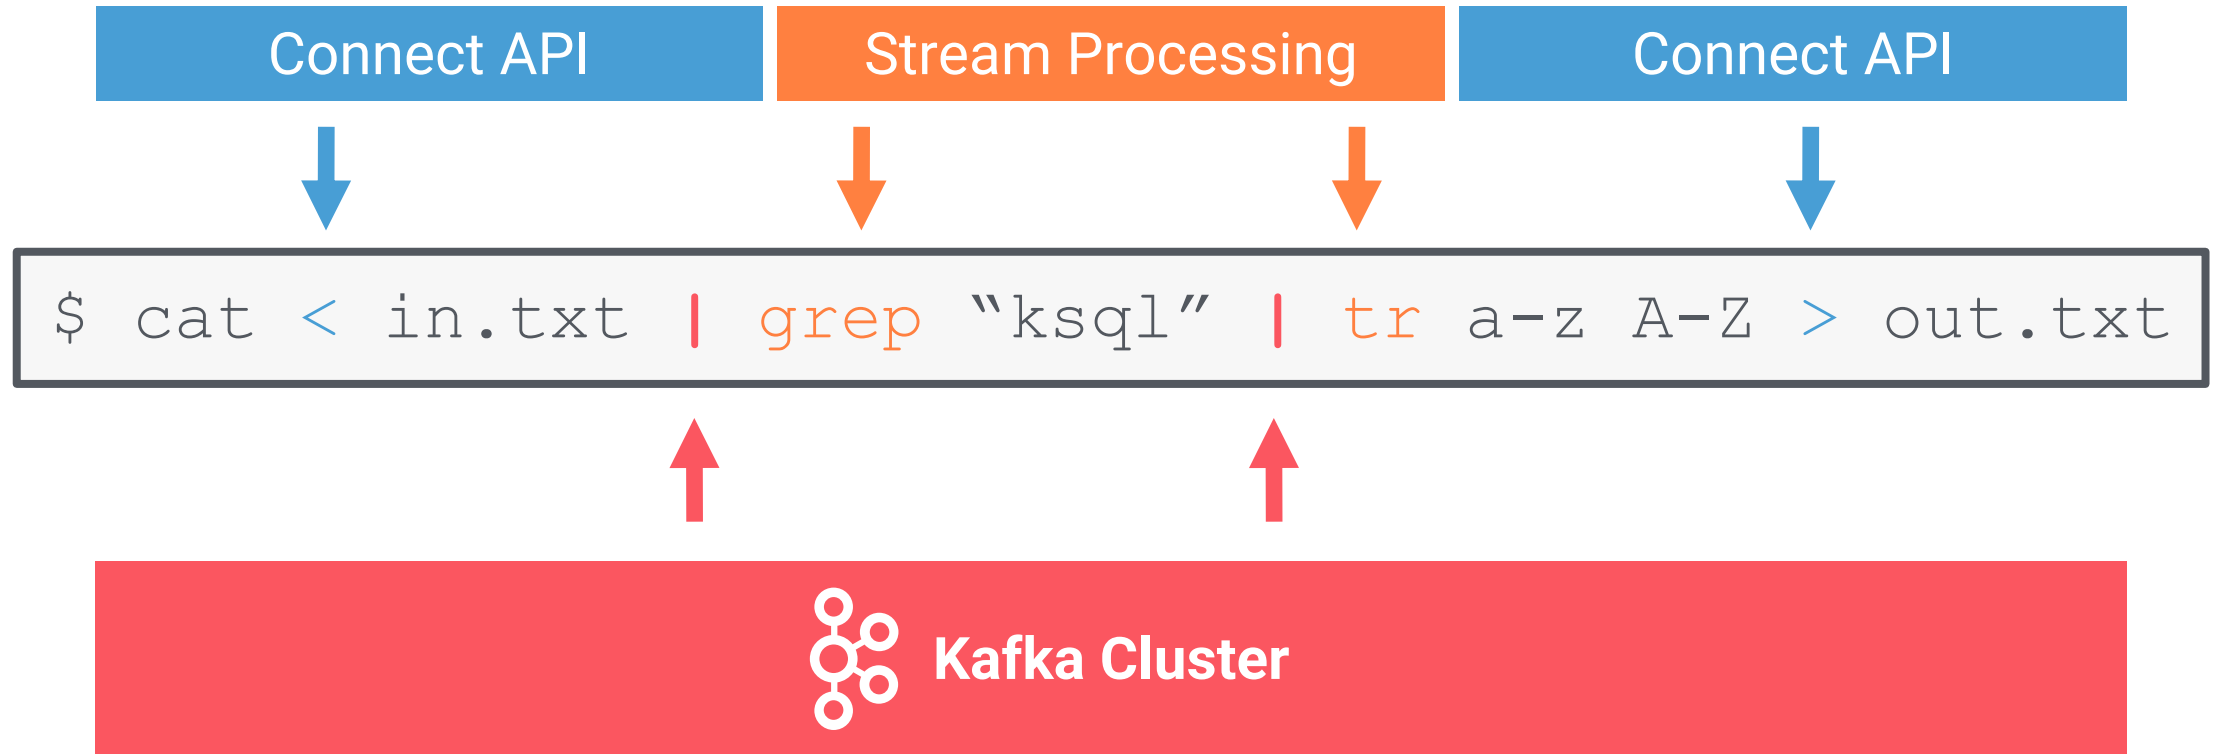
in a durable, fault-
tolerant way.

**Process
streams of data**

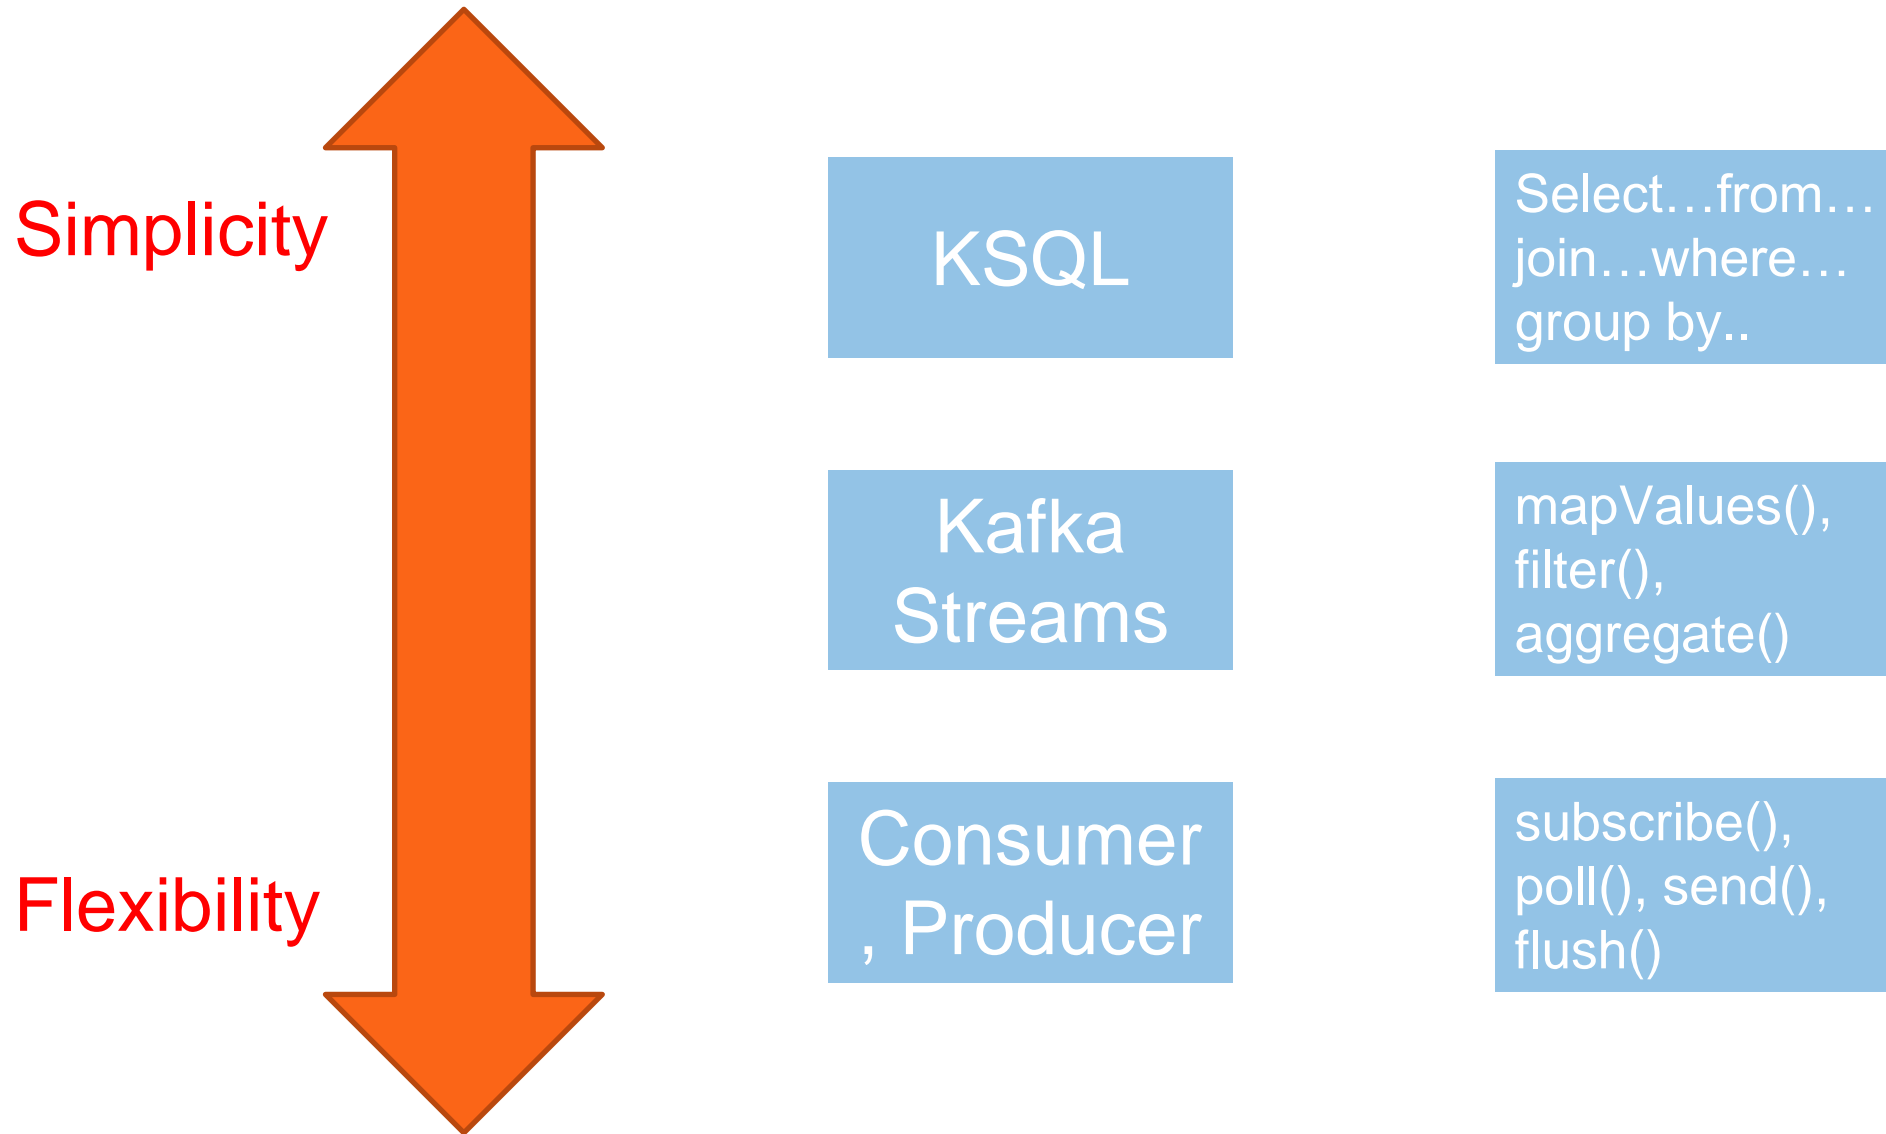


in real time, as they
occur.

Stream Processing by Analogy



Client Trade-offs




Stream processing with Kafka

```
object FraudFilteringApplication extends App {  
  val builder: StreamsBuilder = new StreamsBuilder()
```

```
  val fraudulentPayments: KStream[String, Payment] = builder  
    .stream[String, Payment]("payments-kafka-topic")  
    .filter((_ , payment) => payment.fraudProbability > 0.8)  
  fraudulentPayments.to("fraudulent-payments-topic")
```

```
  val streams: KafkaStreams = new KafkaStreams(builder.build(), config)  
  streams.start()  
}
```


Example: Using **Kafka's Streams API** for writing elastic, scalable, fault-tolerant Java and Scala applications



Stream processing with Kafka

```
CREATE STREAM fraudulent_payments AS  
  SELECT * FROM payments  
  WHERE fraudProbability > 0.8;
```

Same example, now with **KSQL**.
Not a single line of Java or Scala code needed.



Kafka Streams & KSQL - The Easiest Way to Process Data Streams



**Runs
everywhere**




**Clustering
done for you**



**Exactly-once
processing**




**Event-time
processing**



**Integrated
database**



**Joins, windowing,
aggregation**



**S/M/L/XL/XXL/XXXL
sizes**

What is it for ?

Streaming ETL / Enrichment

```
CREATE STREAM clicks_with_city AS  
SELECT c.*, u.city  
FROM clickstream c  
LEFT JOIN users u ON c.user_id = u.id;
```

What is it for ?

Anomaly Detection

- Identifying patterns or anomalies in real-time data, surfaced in milliseconds

```
CREATE TABLE possible_fraud AS  
SELECT card_number, count(*)  
FROM authorization_attempts  
WINDOW TUMBLING (SIZE 5 SECONDS)  
GROUP BY card_number  
HAVING count(*) > 3;
```

What is it for ?

Real Time Monitoring

- Log data monitoring, tracking and alerting
- Sensor / IoT data

```
CREATE TABLE error_counts AS
SELECT error_code, count(*)
FROM monitoring_stream
WINDOW TUMBLING (SIZE 1 MINUTE)
WHERE type = 'ERROR'
GROUP BY error_code;
```

What is it for ?

- Redaction
- Event De-duplication
- Event Re-ordering
- Sessionization
- Validation
- Microservices....

Do you think that's a table you are querying ?



Where is KSQL not such a great fit?

Post-fact Ad-hoc queries

- Limited span of time usually retained in Kafka
- No indexes for random lookups

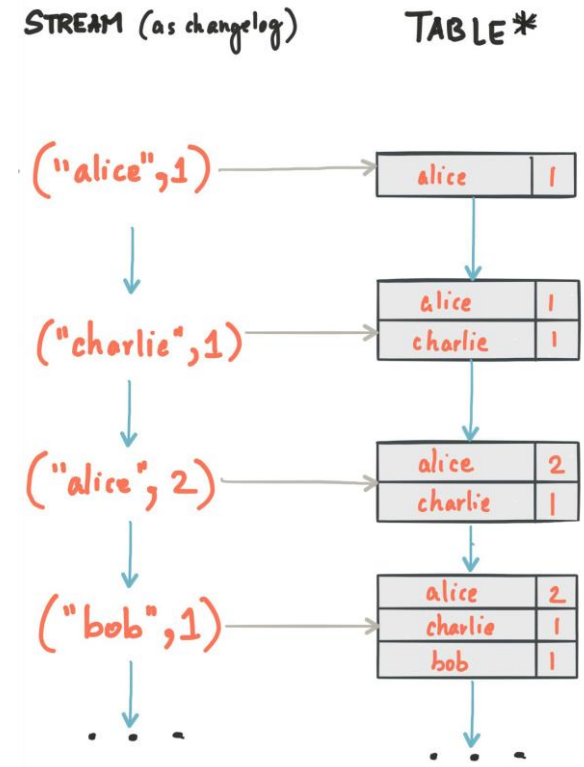
BI reports (Tableau etc.)

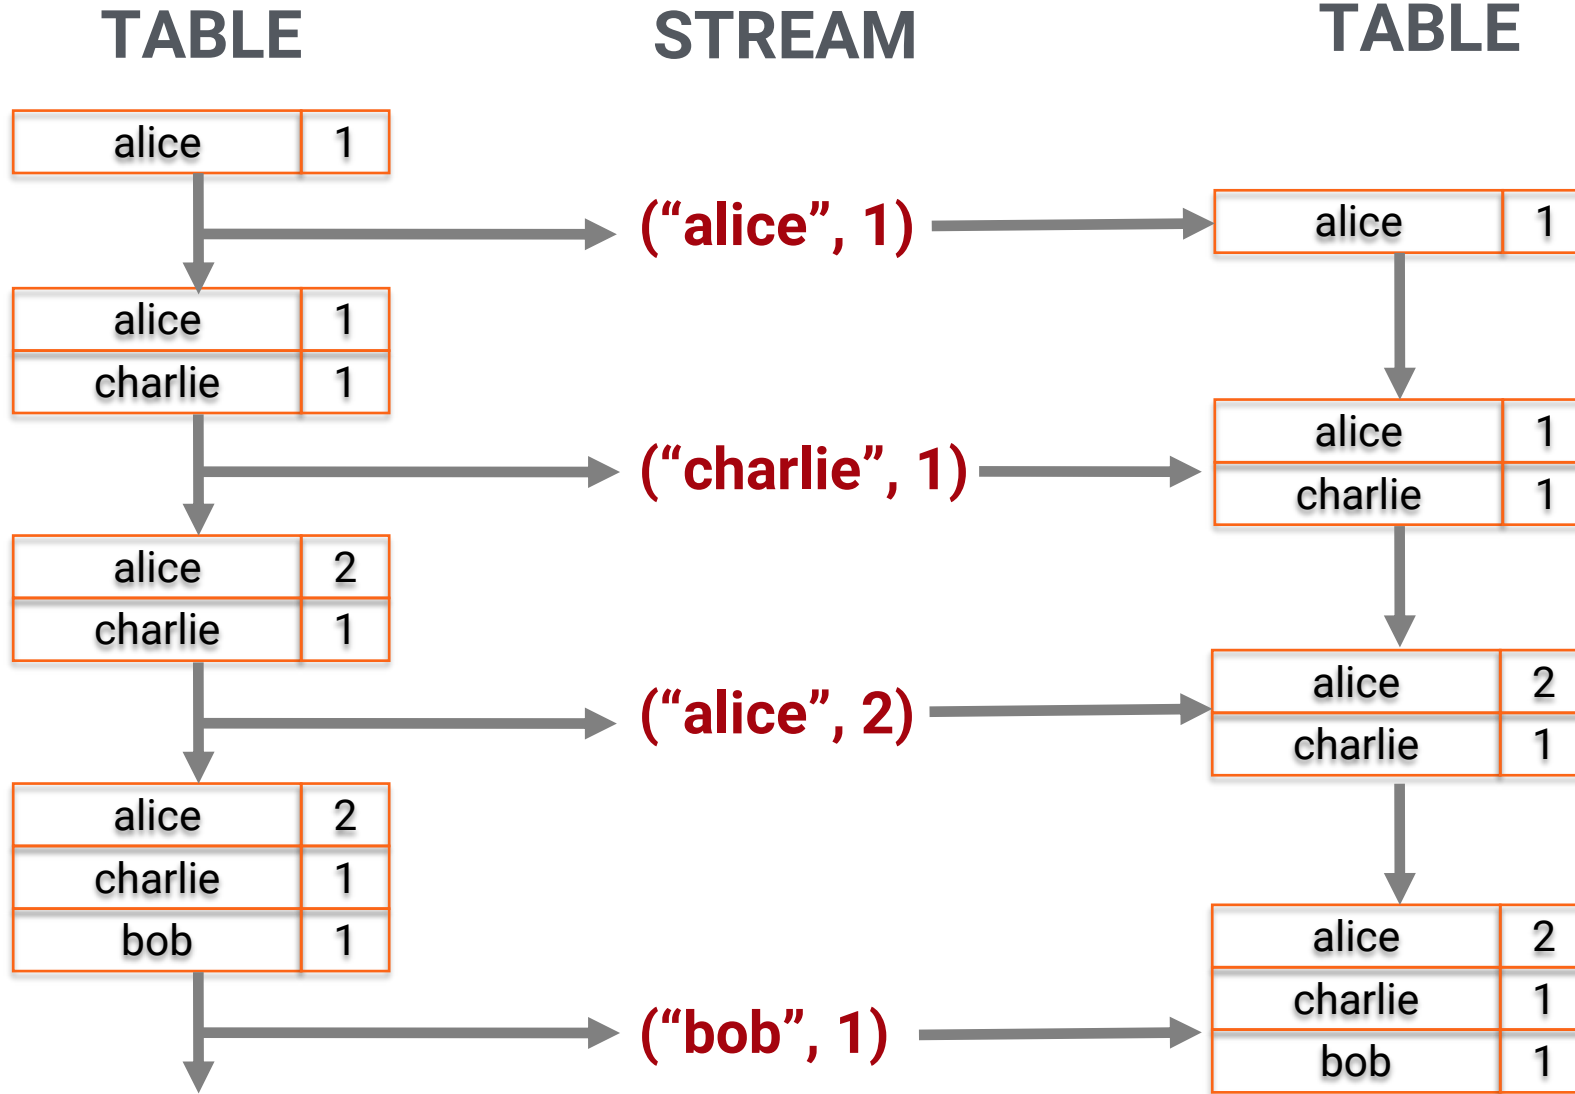
- No secondary indexes
- No JDBC (most BI tools are not good with continuous results!)

Stream/Table Duality

Streams & Tables

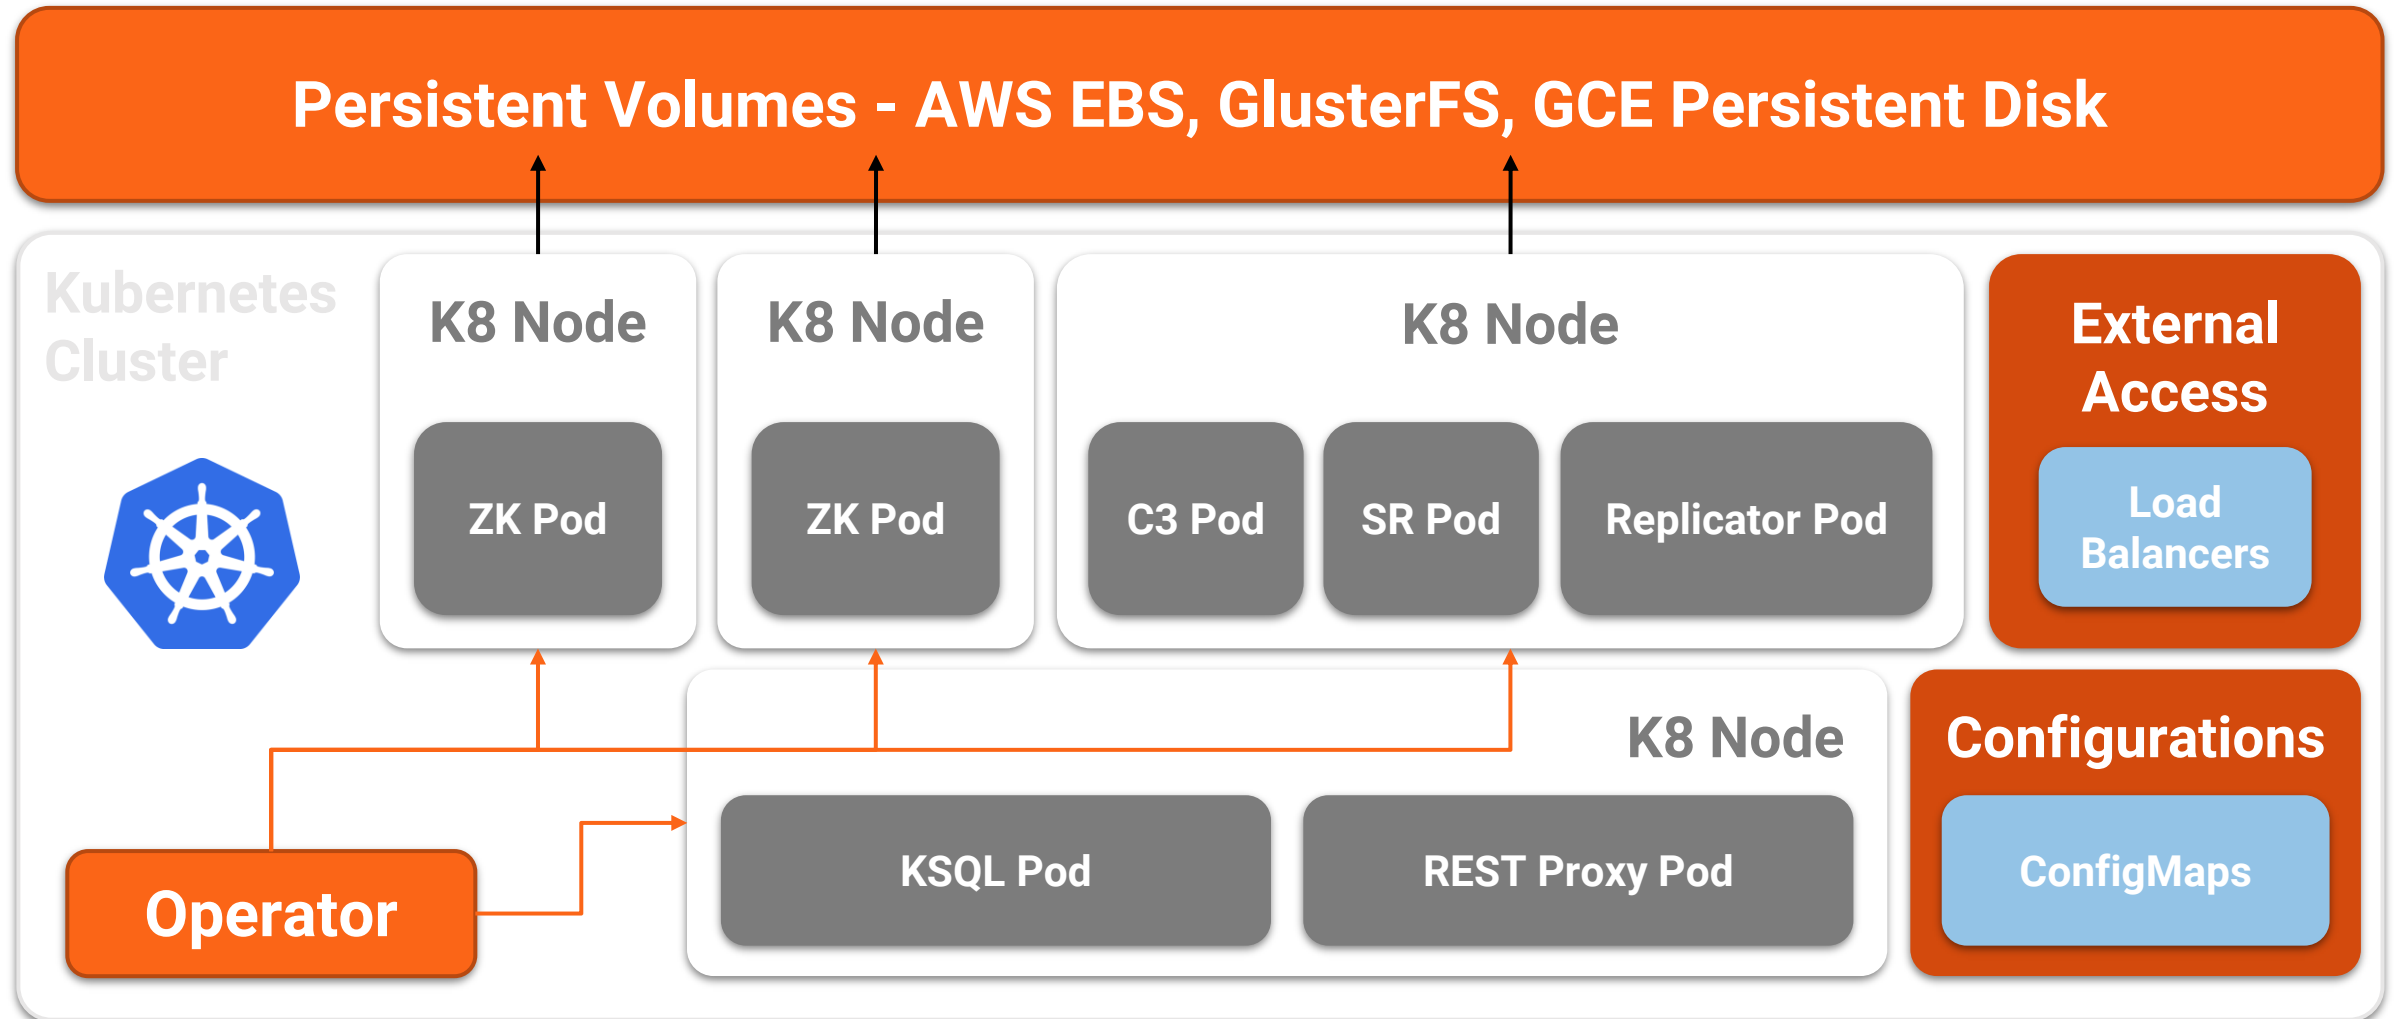
- **STREAM and TABLE as first-class citizens**
- Interpretations of topic content
- **STREAM** - data in motion
- **TABLE** - collected state of a stream
 - One record per key (per window)
 - Current values (compacted topic)
 - Changelog
- **STREAM – TABLE Joins**







Confluent Operator Architecture and Deployment



**[http://ksql-
chicago.gcp.selabs.net](http://ksql-chicago.gcp.selabs.net)**



Learn Kafka.

Start building with Apache Kafka at Confluent Developer.



Confluent Developer
developer.confluent.io

Stay in touch!



Confluent Blog
cnfl.io/blog



Streaming Audio
cnfl.io/podcast



Try Confluent
cnfl.io/download

